


# CS 249r: Special Topics in Edge Computing

Intro to Autonomous Systems / Robotics Wrap-Up

---



Brian Plancher  
Fall 2019





The goal for the next couple of lectures is to develop a **high level** understanding of:

---

1. What is an autonomous system
  2. Key **problems** and **constraints** for autonomous systems
  3. Some of the most important (classes of) **algorithms** in robotics
    - A. The **model based** vs. **model free** tradeoff
    - B. The **online** vs **offline** tradeoff
    - C. The **no free lunch** theorem and the need for **approximations**
  4. How **computer systems / architecture** design has and can play a role in improving autonomous systems
-

# What do we mean by an Autonomous System?




The goal for the next couple of lectures is to develop a **high level** understanding of:

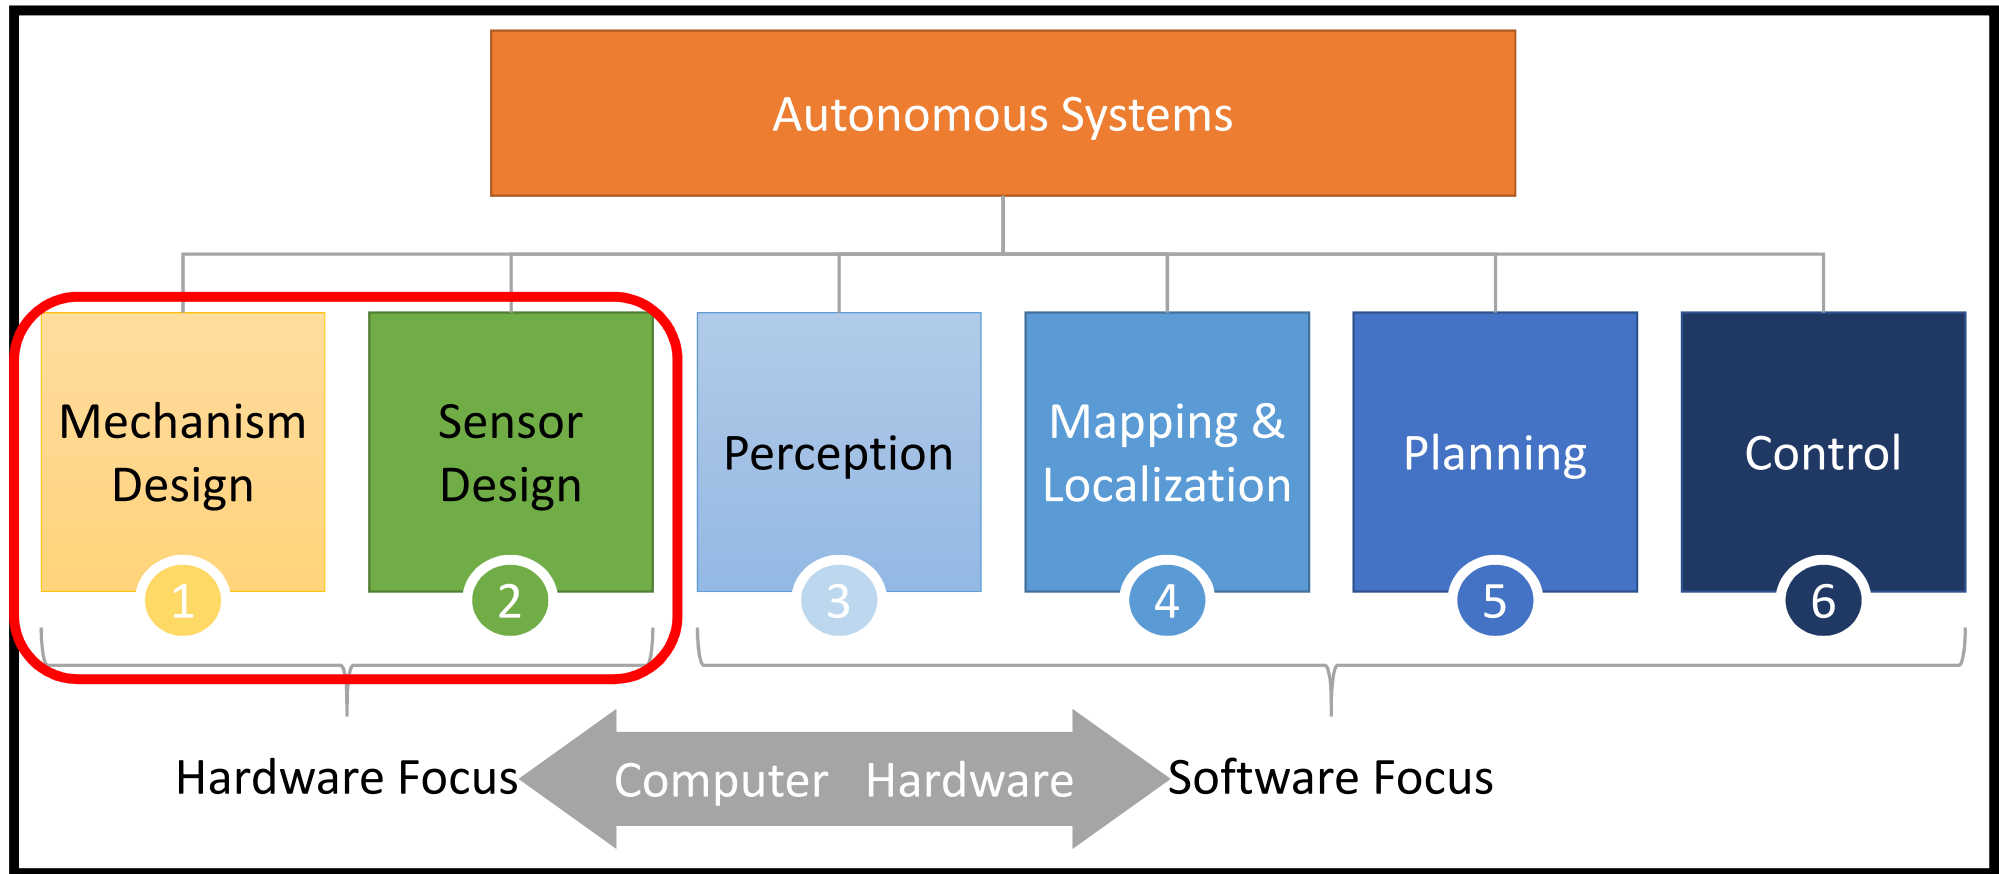
---

1. What is an autonomous system
  2. Key **problems** and **constraints** for autonomous systems
  3. Some of the most important (classes of) **algorithms** in robotics
    - A. The **model based** vs. **model free** tradeoff
    - B. The **online** vs **offline** tradeoff
    - C. The **no free lunch** theorem and the need for **approximations**
  4. How **computer systems / architecture** design has and can play a role in improving autonomous systems
-

# Autonomous Systems / Robotics is a BIG space



# Autonomous Systems / Robotics is a BIG space



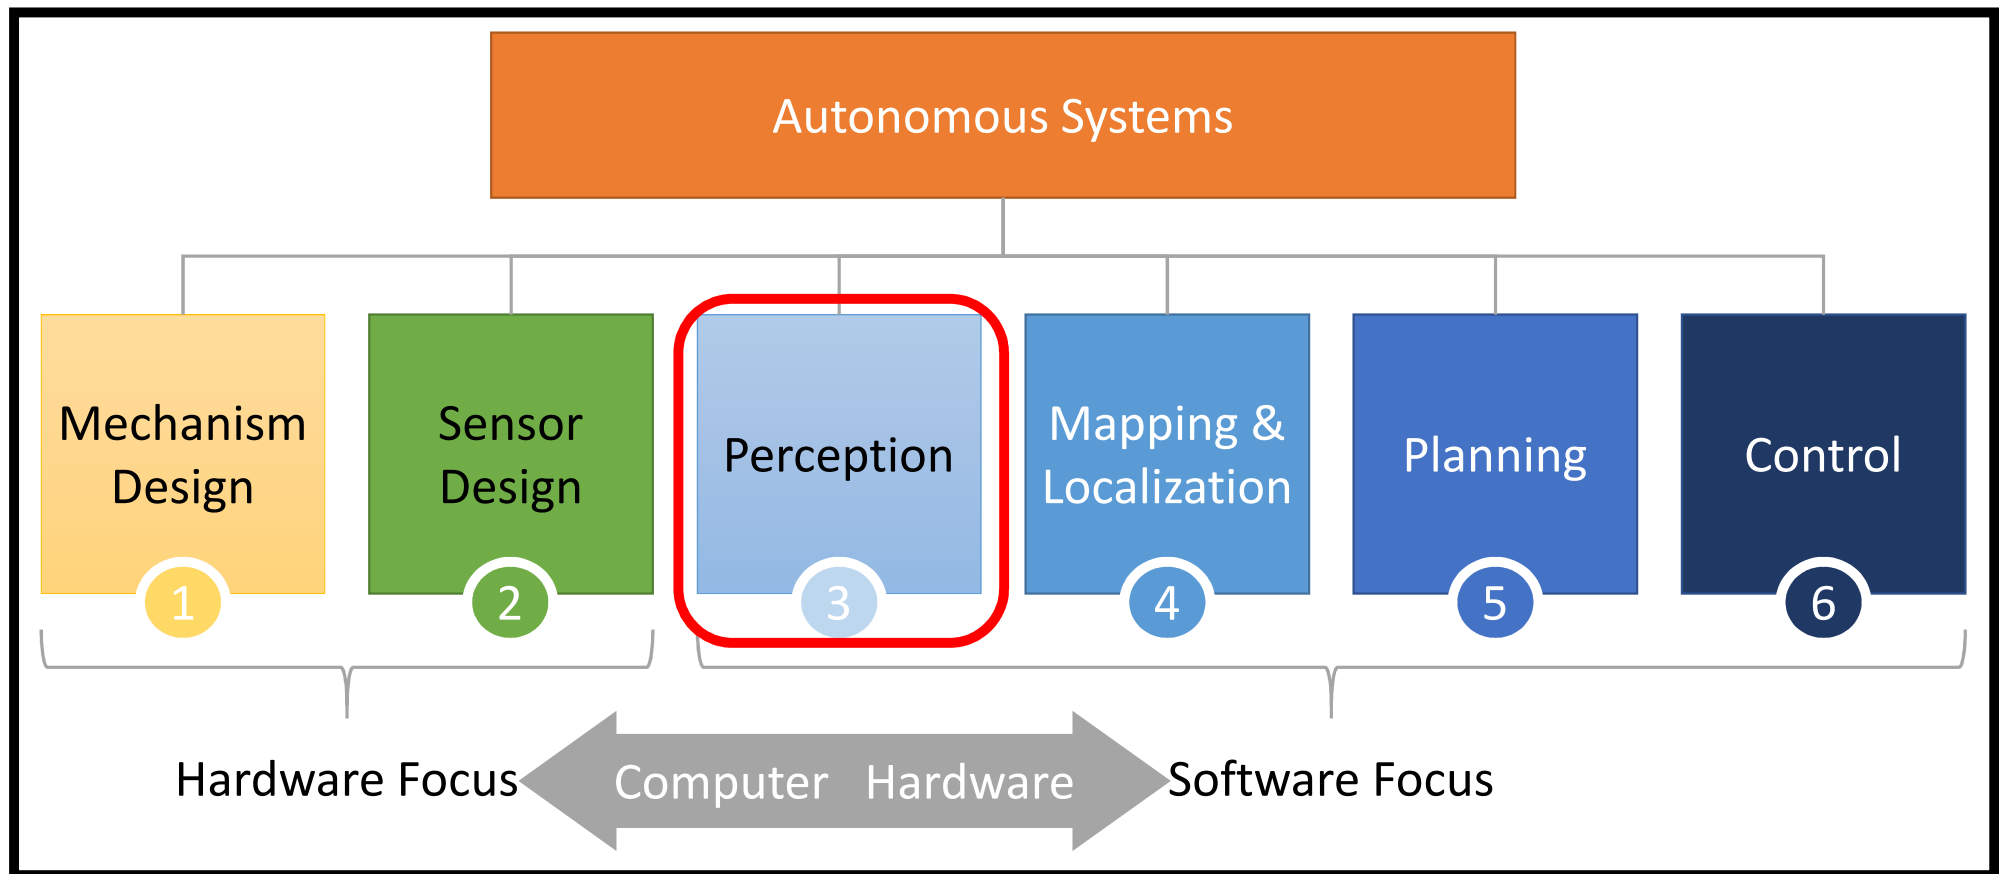
## 1 2 Key Takeaways:

---



1. When designing algorithms for robots you need to understand the physical capabilities of the robot and you (potentially) need to understand how to model its physical behaviors
  2. Different kinds of systems will have different power, weight, and performance budgets for computer hardware
-

# Autonomous Systems / Robotics is a BIG space






3

# Computer Vision (and Perception in general) is hard

Retinal color

$$\mathbf{c}(\ell(\lambda)) = (c_s, c_m, c_l)$$





Perceived color

Object color

Color names

$$c_s = \int k_s(\lambda) \ell(\lambda) d\lambda$$




3

CV/Perception is solved by modeling and approximating the classification of convolution

Retinal color

$$\mathbf{c}(\ell(\lambda)) = (c_s, c_m, c_l)$$





Perceived color

Object color

Color names

**“Classification”**

$$c_s = \int k_s(\lambda) \ell(\lambda) d\lambda$$



**“Convolution”**




3

We approximate convolution using linear filters


$$G_{\sigma} = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$

|       |       |       |       |       |
|-------|-------|-------|-------|-------|
| 0.003 | 0.013 | 0.022 | 0.013 | 0.003 |
| 0.013 | 0.059 | 0.097 | 0.059 | 0.013 |
| 0.022 | 0.097 | 0.159 | 0.097 | 0.022 |
| 0.013 | 0.059 | 0.097 | 0.059 | 0.013 |
| 0.003 | 0.013 | 0.022 | 0.013 | 0.003 |


5 x 5,  $\sigma = 1$



No smoothing



$\sigma = 2$





$\sigma = 4$

3

Deep learning automates the design of filters, and the selection/combination of features for classification

### AlexNet: the first widely successful application of deep learning




<https://papers.nips.cc/paper/4824-imagenet-classification-with-deep-convolutional-neural-networks.pdf>

<https://cv-tricks.com/cnn/understand-resnet-alexnet-vgg-inception/>

3

Deep learning automates the design of filters, and the selection/combination of features for classification

### AlexNet: the first widely successful application of deep learning




<https://papers.nips.cc/paper/4824-imagenet-classification-with-deep-convolutional-neural-networks.pdf>

<https://cv-tricks.com/cnn/understand-resnet-alexnet-vgg-inception/>

3

Deep learning automates the design of filters, and the selection/combination of features for classification

### AlexNet: the first widely successful application of deep learning



<https://papers.nips.cc/paper/4824-imagenet-classification-with-deep-convolutional-neural-networks.pdf>

<https://cv-tricks.com/cnn/understand-resnet-alexnet-vgg-inception/>



3

Deep learning automates the design of filters, and the selection/combination of features for classification


But watch out for adversarial attacks on the math!



“panda”

57.7% confidence

+  $\epsilon$



=

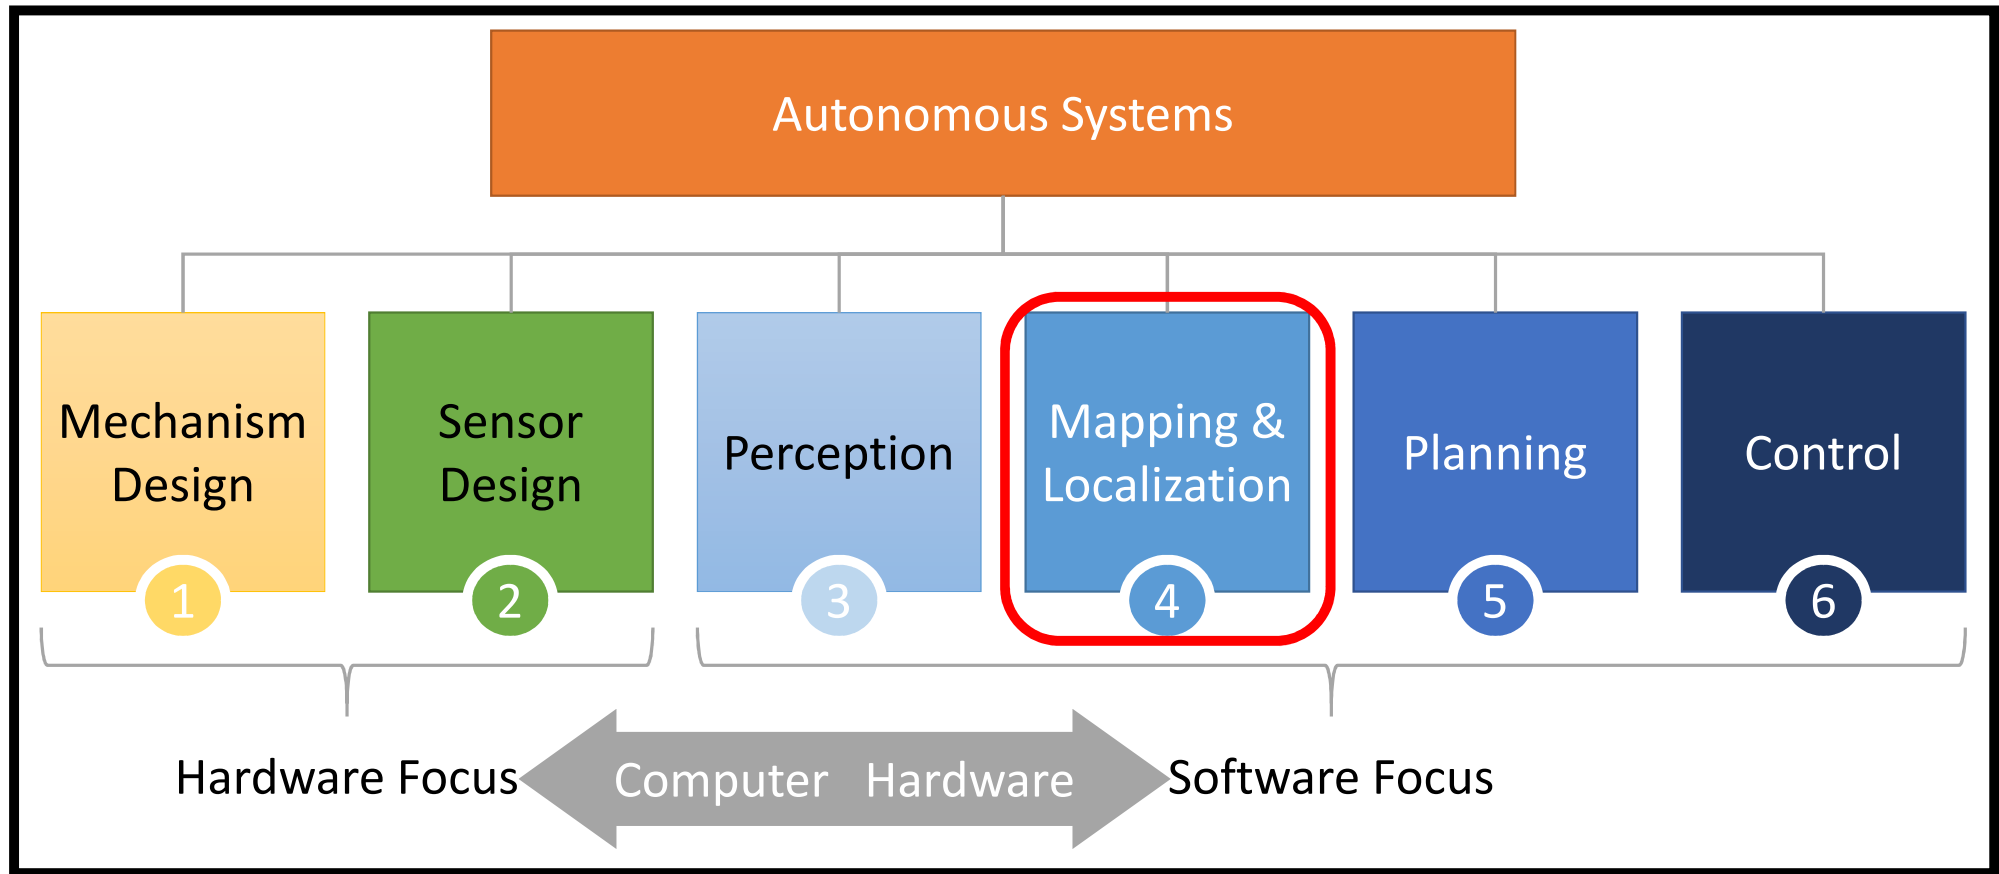


“gibbon”

99.3% confidence


“No Free Lunch!”

# Autonomous Systems / Robotics is a BIG space







## 4 Mapping/Localization is hard



### Three Problems

1. GPS is only accurate to  $O(10m)$
2. GPS relies on already having a perfect map of the environment (unrealistic often)
3. Other sensor data is also quite noisy!



4


Mapping/Localization is solved by modeling the world as an HMM and using modeling and approximating to solve it

Track the **Belief State**  $B_t$  of the state and landmarks


$$B_t = p(X_t | X_0, E_0 \dots E_{t-1})$$

**Hidden Markov Model (HMM)**

States  $X$  update in time but we only observe the effects  $E$



**Time Update**



$$P(x_{t+1} | x_0, e_1 \dots e_t) = \int P(x_t | x_0, e_1 \dots e_t) * P(x_{t+1} | x_t)$$

**Evidence Update**




$$P(x_{t+1} | x_0, e_1 \dots e_{t+1}) \propto \int P(x_{t+1} | x_0, e_1 \dots e_t) * P(e_{t+1} | x_{t+1})$$

4

Mapping/Localization is solved by modeling the world as an HMM and using modeling and approximating to solve it

Time Update




$$P(x_{t+1}|x_0, e_1 \dots e_t) = \int P(x_t|x_0, e_1 \dots e_t) * P(x_{t+1}|x_t)$$

Evidence Update



$$P(x_{t+1}|x_0, e_1 \dots e_{t+1}) \propto \int P(x_{t+1}|x_0, e_1 \dots e_t) * P(e_{t+1}|x_{t+1})$$

Particle Filter




Approximate  
with Samples

Model  
With  
Gaussian


Approximate  
as Linear

Approximate  
with  
Samples

Extended Kalman Filter (EKF)




Unscented Kalman Filter - UKF



4

Mapping/Localization is solved by modeling the world as an HMM and using modeling and approximating to solve it

Time Update




$$P(x_{t+1}|x_0, e_1 \dots e_t) = \int P(x_t|x_0, e_1 \dots e_t) * P(x_{t+1}|x_t)$$

Evidence Update




$$P(x_{t+1}|x_0, e_1 \dots e_{t+1}) \propto \int P(x_{t+1}|x_0, e_1 \dots e_t) * P(e_{t+1}|x_{t+1})$$

Particle Filter


Approximate  
with SamplesModel  
With  
GaussianApproximate  
near

Models / Approx  
= Some Error  
"No Free Lunch!"

Extended Kalman Filter (EKF)

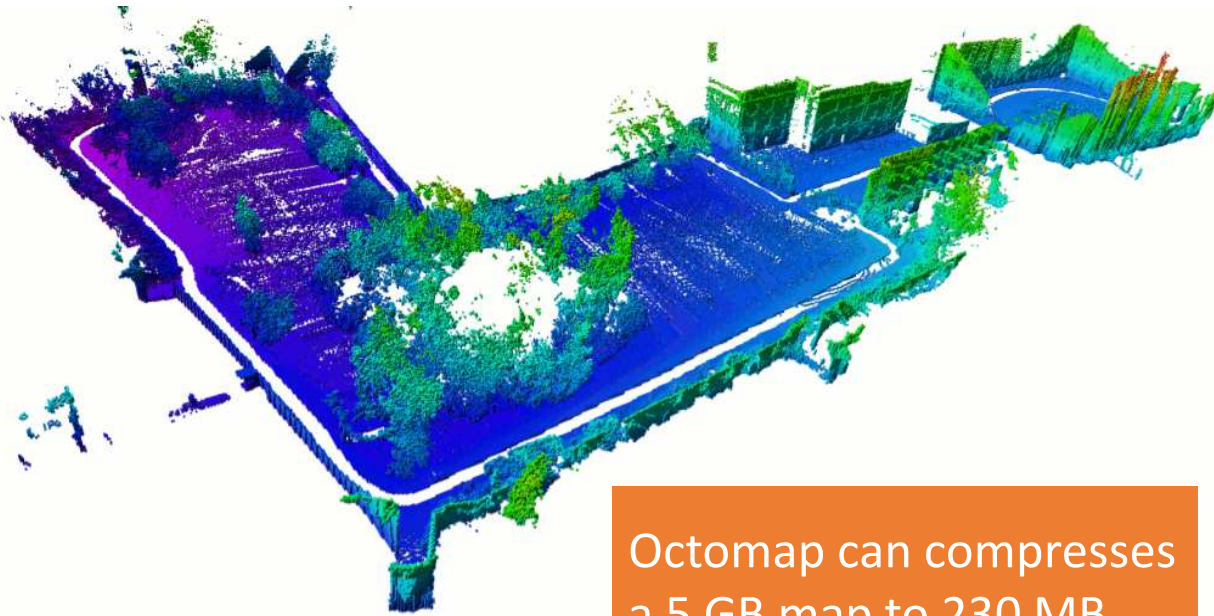


Unscented Kalman Filter - UKF




4

Also we need to approximate the resolution of our maps and store them intelligently to fit them in memory



Octomap can compresses  
a 5 GB map to 230 MB

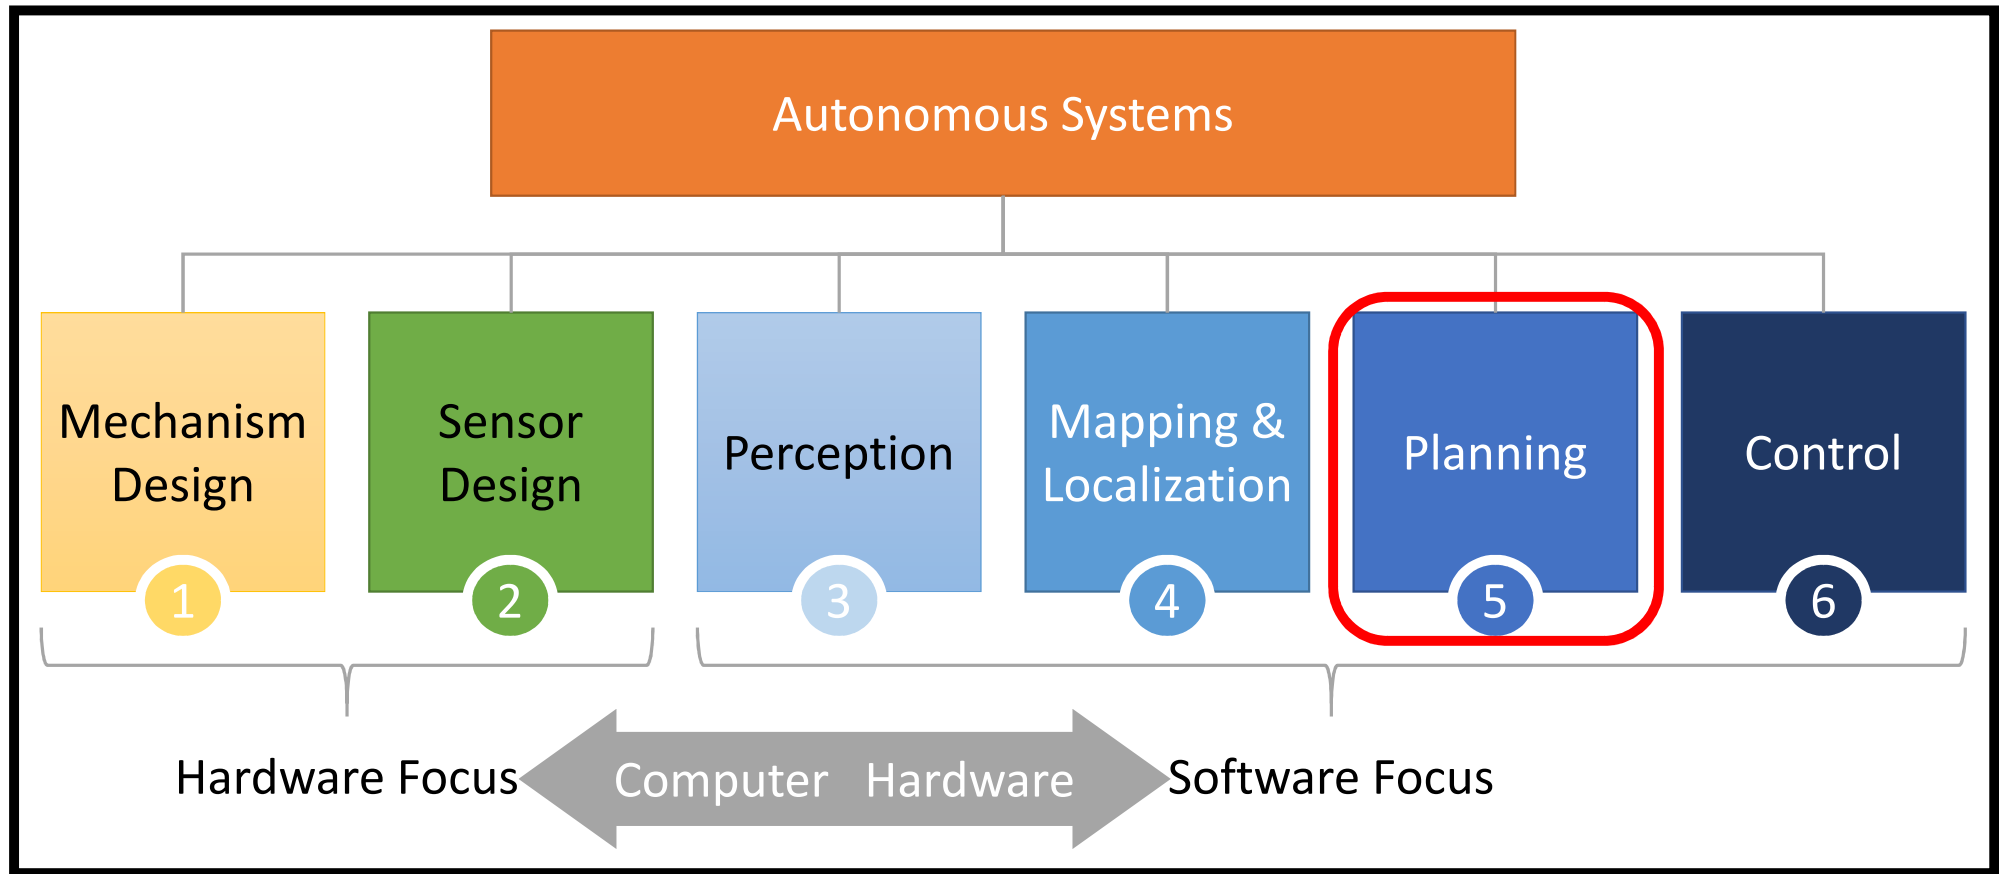


**Fig. 2** Example of an octree storing free (shaded white) and occupied (black) cells. The volumetric model is shown on the left and the corresponding tree representation on the right.




**Fig. 3** By limiting the depth of a query, multiple resolutions of the same map can be obtained at any time. Occupied voxels are displayed in resolutions 0.08 m, 0.64, and 1.28 m.

# Autonomous Systems / Robotics is a BIG space





## 5 Planning (in Configuration Space) is hard





One approach is to discretize the statespace (grid it) and use graph search ( $A^*$  = fast)

Another is to solve a global optimization problem:


$$\begin{aligned} & \underset{s_0, a_0, \dots, s_N, a_N}{\text{minimize}} \sum_{k=0}^N c(s_k, a_k) \\ & \text{subject to } s_{k+1} = f(s_k, a_k) \\ & \quad s_N = s_{\text{goal}} \end{aligned}$$

Complexity scales with  $d^{|S|=|A|}$ : **Curse of Dimensionality**



5

There are three main ways to approximately plan in Configuration Space




Random Search

### Machine Learning



Humanoid:  
27 DoFs, 21 Actuators.



Local Search



5

We can approximately plan locally optimal plans in Configuration Space in three ways

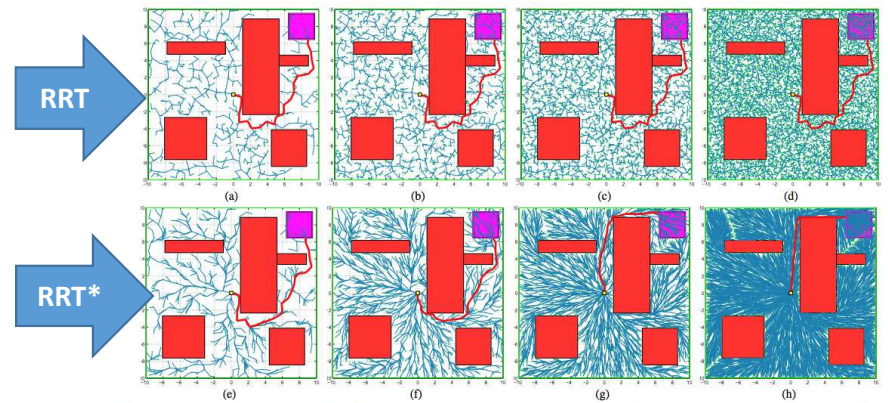






Fig. 1. A Comparison of the RRT\* and RRT algorithms on a simulation example. The tree maintained by the RRT algorithm is shown in (a)-(d) in different stages, whereas that maintained by the RRT\* algorithm is shown in (e)-(h). The tree snapshots (a), (e) are at 1000 iterations, (b), (f) at 2500 iterations, (c), (g) at 5000 iterations, and (d), (h) at 15,000 iterations. The goal regions are shown in magenta. The best paths that reach the target are highlighted with red.

The PRM is searched for a path from  $s$  to  $g$



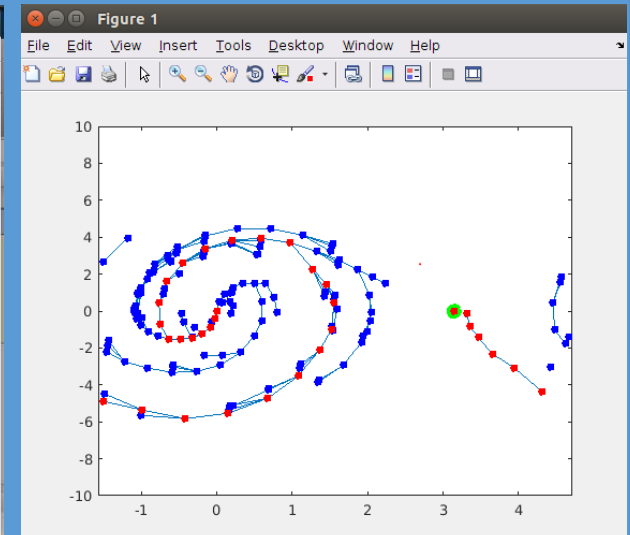
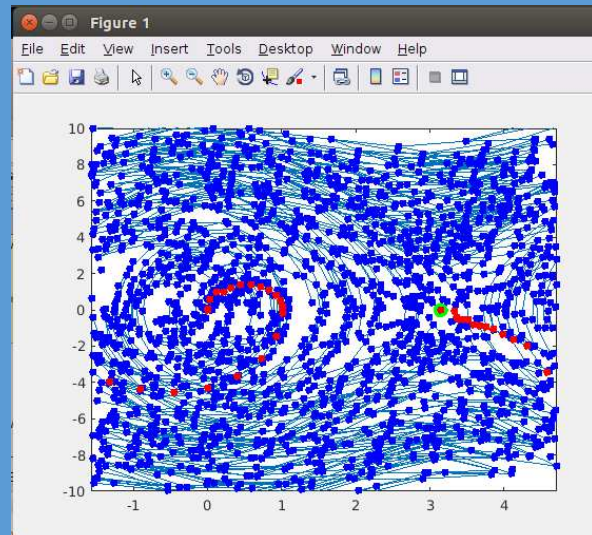
5

We can approximately plan locally optimal plans in Configuration Space in three ways



Random Search

Note: Can scale to low-dimensional dynamical systems with LQR-RRT\*



"No Free Lunch!"


The PRM is searched for a path from s to g

(g) at 5000 iterations, and (d), (h) at 15,000 iterations. The goal regions are shown in magenta. The best paths that reach the target are highlighted with red.

5

We can approximately plan locally optimal plans in Configuration Space in three ways

## Machine Learning




**My two cents:  
Yes, and no free  
lunch!**

**Needs to re-learn  
physics and suffers  
from sample  
complexity**

**In two weeks more  
on this!**



5

We can approximately plan locally optimal plans in Configuration Space in three ways



Local Search

Solve math locally with linear & quadratic approximations




5

## Practical Challenges for Trajectory Optimization: Not Complete, Not Robustness and Contact = No Free Lunch!

1. **Not complete** (aka no guaranteed solution) and often slow!
2. Solvers are **numerically sensitive**
3. Solutions are sensitive to initial trajectories and **perturbations**
4. The physics equations are fundamentally different when an object makes or breaks **contact** leading to a **combinatorial explosion**

One approach to avoid solving these large hard problems is to solve the problem by **combining simpler models** of the system although this leads to **conservative** behavior





5 Control is hard (even for the experts)

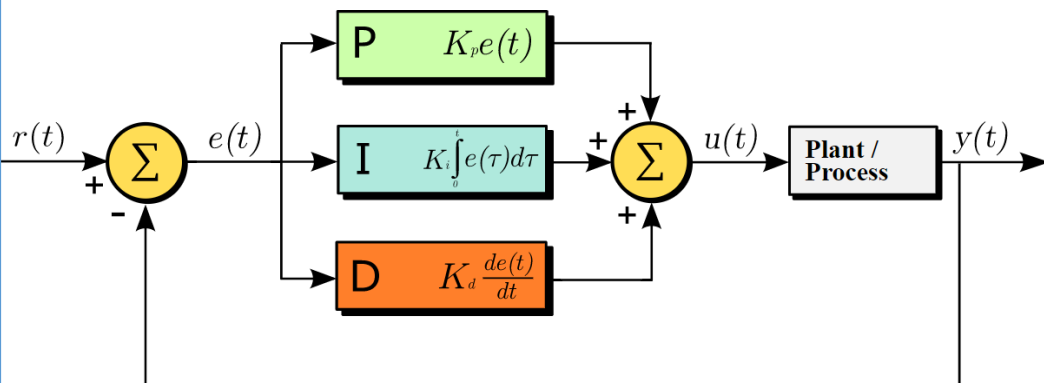


6

We use feedback tracking controllers to run our plans in the real world (and handle the differences encountered)

Model as linear combination of errors and approximate gains

This is the canonical PID controller!



LQR: Quadratic Cost with Linear Dynamics

$$\min_{x,u} \sum_{k=0}^N (x_k - x_g)^T Q (x_k - x_g) + u_k^T R u_k$$

$$\text{s.t. } x_{k+1} = A x_k + B u_k$$

$$u_k = -K_k x_k$$

Solve math locally with linear & quadratic approximations

6



We use feedback tracking controllers to run our plans in the real world (and handle the differences encountered)

Model as linear combination of basis functions and approximate gains

And if we can plan fast enough we just use constant replanning to control (MPC)

Cost with dynamics

This is the cost



We'll see this again next Wednesday!

$$J = \sum_{k=0}^{N-1} (x_k - x_g)^T R u_k + u_k^T R u_k + B u_k$$



## 6 Practical Challenges for Control: Contact




Figure 17.1 - Modeling contact as a hybrid system.

The goal for the next couple of lectures is to develop a **high level** understanding of:

---

1. What is an autonomous system
  2. Key **problems** and **constraints** for autonomous systems
  3. Some of the most important (classes of) **algorithms** in robotics
    - A. The **model based** vs. **model free** tradeoff
    - B. The **online** vs **offline**
    - C. The **no free lunch** t
  4. How **computer systems / architecture** design has and can play a role in improving autonomous systems
- 

This is what we will explore

in all of the papers! **optimizations**

The goal for the next couple of lectures is to develop a **high level** understanding of:

1. What is an autonomous system?


2. Key **problems** and **concepts**

3. Some of the most important

A. The **model based** vs. **model free**

B. The **online** vs **offline**

C. The **no free lunch** theorem



This is what we will explore  
in all of the papers!

applications

4. How **computer systems / architecture** design has and can play a role in improving autonomous systems

Your homework – get on HOTCRP

---

Email **Glenn Holloway:**  
[holloway@eecs.harvard.edu](mailto:holloway@eecs.harvard.edu)

He will send you a password (username is that email address) after which I can assign you access to review papers

---